

13-simple-linear-regression

May 17, 2023

1 Simple Linear Regression Model

```
[3]: #Step 1: Import the necessary libraries

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from sklearn import datasets, linear_model, metrics
```

```
[7]: # Step 2: # load the boston dataset

#boston = datasets.load_boston(return_X_y=False)

boston = pd.read_csv('C:/Users/lenovo/Desktop/DS Lab Week Wise/
↳Week-5-Regression_model/Boston.csv')
```

```
[8]: # Displaying the dataset

boston.head()
```

```
[8]:
```

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	\
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	
2	3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	
3	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	
4	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	

	tax	ptratio	black	lstat	medv
0	296	15.3	396.90	4.98	24.0
1	242	17.8	396.90	9.14	21.6
2	242	17.8	392.83	4.03	34.7
3	222	18.7	394.63	2.94	33.4
4	222	18.7	396.90	5.33	36.2

```
[11]: # Step 3: Having the glance at independent and dependent variable

boston = boston.loc[:,['lstat','medv']]
```


```
[13]: # Displaying the new dataset
```

```
boston.head(10)
```

```
[13]:   lstat  medv
0    4.98  24.0
1    9.14  21.6
2    4.03  34.7
3    2.94  33.4
4    5.33  36.2
5    5.21  28.7
6   12.43  22.9
7   19.15  27.1
8   29.93  16.5
9   17.10  18.9
```

```
[14]: # Step 4: Visualizing the change in the variables
```

```
boston.plot(x = 'lstat', y = 'medv', style = 'o')
plt.xlabel('lstat')
plt.ylabel('medv')
plt.show()
```



```
[16]: # Step 5: Dividing the dataset into dependent and independent variables
```

```
x = boston['lstat']  
y = boston['medv']  
  
x
```

```
[16]: 0      4.98  
      1      9.14  
      2      4.03  
      3      2.94  
      4      5.33  
      ...  
     501      9.67  
     502      9.08  
     503      5.64  
     504      6.48  
     505      7.88  
      Name: lstat, Length: 506, dtype: float64
```

```
[17]: # Step 6: Splitting the data into train and test sets
```

```
from sklearn.model_selection import train_test_split  
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.  
    ↪4, random_state=1)
```

```
[18]: # Checking the shapes of the train and test sets
```

```
print(x_train.shape)  
print(x_test.shape)  
print(y_train.shape)  
print(y_test.shape)
```

```
(303,)  
(203,)  
(303,)  
(203,)
```

```
[30]: x_train = x_train[:, None]
```

```
x_test = x_test[:, None]
```

C:\Users\lenovo\AppData\Local\Temp\ipykernel_6060\296195278.py:3: FutureWarning: Support for multi-dimensional indexing (e.g. `obj[:, None]`) is deprecated and will be removed in a future version. Convert to a numpy array before indexing instead.

```
x_test = x_test[:, None]
```

```
[23]: # Step 7: Training the algorithm

from sklearn.linear_model import LinearRegression

lregressor = LinearRegression()
lregressor.fit(x_train.reshape(-1, 1), y_train)
```

```
[23]: LinearRegression()
```

```
[31]: # Step 8: Predicting the y values from the built model

y_pred = lregressor.predict(x_test)
```

```
[24]: # Step 9: Retrieving the intercept

print(lregressor.intercept_)
```

```
34.13076276013123
```

```
[25]: # Step 10: Retrieving the slope

print(lregressor.coef_)
```

```
[-0.91573274]
```

```
[32]: # Step 11: Printing the predicted values

y_pred
```

```
[32]: array([27.23529525, 27.55580171, 16.91498731, 26.71332759, 24.79944617,
23.94781473, 29.8268189 , 22.18960787, 17.72083211, 26.06315735,
27.04299138, 29.90923484, 21.66764021, 24.75365953, 23.39837508,
23.02292466, 12.84913395, 29.89092019, 27.32686852, 7.08917504,
23.59983629, 18.88381269, 25.63276296, 28.52647841, 29.85429088,
11.79604131, 15.50475889, 24.48809704, 27.48254309, 15.0377352 ,
29.10339003, 17.22633644, 31.49345248, 19.07611657, 25.80675218,
21.68595487, 17.83987737, 29.24990727, 12.75756068, 20.41308636,
27.40928447, 27.94956679, 27.17119396, 12.0616038 , 17.62010151,
13.30700032, 32.37255591, 19.15853251, 25.14742461, 24.39652377,
23.45331905, 23.85624145, 29.3597952 , 23.96612938, 6.95181513,
27.96788144, 6.68625263, 28.65468099, 20.69696351, 30.51361845,
20.37645705, 28.15102799, 15.88936664, 17.97723728, 7.07086038,
29.51546977, 31.8689029 , 26.20967459, 24.6071423 , 23.4716337 ,
28.49900643, 8.23384096, 21.31050444, 23.16028457, 21.2280885 ,
24.58882764, 31.42019386, 26.59428234, 27.36349783, 30.87075422,
20.03763594, 24.69871557, 29.19496331, 12.82166197, 28.65468099,
29.57957106, 16.44796361, 29.47884046, 21.21893117, 18.59993554,
```

```

28.96603012, 29.44221115, 17.5834722 , 22.74820484, 20.83432342,
21.97898934, 25.27562719, 26.65838363, 30.33962923, 22.22623718,
19.76291612, 16.86920067, 28.81035556, 27.84883618, 5.91703714,
22.83062079, 18.41678899, 29.09423271, 22.08887727, 27.10709267,
31.23704731, 25.85253882, 14.42419426, 28.20597195, 26.99720474,
31.50260981, 25.2115259 , 23.42584707, 32.31761194, 27.3177112 ,
25.00090737, 16.99740325, 29.69861631, 21.13651522, 25.42214443,
22.2994958 , 26.27377588, 25.81590951, 20.67864885, 29.28653658,
21.17314453, 24.64377161, 28.44406246, 10.61474608, 27.5191724 ,
27.78473489, 9.57081076, 16.72268343, 27.20782327, 18.37100236,
29.03013142, 17.22633644, 17.04318989, 28.30670255, 31.08137275,
25.74265089, 27.53748705, 14.60734081, 28.04114006, 17.71167479,
25.34888581, 15.92599595, 26.56681035, 20.90758204, 14.68059943,
30.17479734, 27.61074567, 27.0521487 , 20.70612084, 25.39467245,
30.86159689, 29.28653658, 28.12355601, 19.75375879, 21.21893117,
22.29033847, 28.30670255, 12.43705422, 18.38931701, 27.97703877,
24.7078729 , 17.59262953, 18.30690107, 28.91108616, 20.61454756,
27.22613792, 28.01366808, 23.11449793, 32.54654513, 26.081472 ,
14.47913823, 26.62175432, 14.69891408, 26.2463039 , 11.64036674,
20.56876093, 24.83607548, 28.10524135, 25.68770693, 20.97168333,
19.74460146, 17.52852824, 22.22623718, 21.82331478, 26.91478879,
20.66949153, 29.9275495 , 25.62360563, 30.05575208, 25.99905606,
18.93875665, 2.65702859, 25.54118969])

```

```
[27]: # Step 12: Printing the Actual values
```

```
y_test
```

```

[27]: 307    28.2
      343    23.9
      47    16.6
      67    22.0
      362   20.8
      ...
      186   50.0
      372   50.0
      442   18.4
      412   17.9
      213   28.1
      Name: medv, Length: 203, dtype: float64

```

```
[33]: # Step 13: Evaluating the algorithm
```

```

from sklearn import metrics
print('Mean Absolute Error: ', metrics.mean_absolute_error(y_test,y_pred))
print('Mean Squared Error: ', metrics.mean_squared_error(y_test,y_pred))

```

```
print('Root Mean Squared Error: ', np.sqrt(metrics.  
↪mean_squared_error(y_test,y_pred)))
```

Mean Absolute Error: 4.85248058295052

Mean Squared Error: 45.23188062977258

Root Mean Squared Error: 6.7254650865031325

[]: